

# TableGAN Implementation

## Introduction:

TableGAN is a specialized variant of Generative Adversarial Networks (GANs) tailored for generating synthetic tabular data. This implementation is designed to mimic the statistical properties of real datasets while ensuring privacy, making it suitable for tasks where data sharing is restricted due to confidentiality concerns. The architecture comprises three main components: a generator, a discriminator, and a classifier, each integral to the model's ability to learn and generate high-quality synthetic data.

## Detailed Architecture

### 1. Generator

- Purpose: Generates synthetic data that mimics the distribution of real data.
- Input: Noise vector of size  $z\_dim$ .
- Output: Synthetic data resembling reshaped tabular data with dimensions (32, 32, 1).

#### Key Layers:

- Dense Layer: Expands the input noise vector into a higher-dimensional tensor.
- Batch Normalization and LeakyReLU: These layers aid in stabilizing the learning process and introducing non-linearity.
- Conv2DTranspose Layers: Used to upsample the data to the desired output shape, facilitating the generation of 2D data structures from flattened inputs.
- Activation: The tanh function normalizes the output, ensuring that the data distribution closely matches that of the input domain.

### 2. Discriminator

- Purpose: Differentiates between real data from the dataset and fake data generated by the generator.
- Input: Data of shape (32, 32, 1).
- Output: A scalar value indicating the authenticity of the input data and a feature vector representing the input.

#### Key Layers:

- Conv2D Layers: Extract features from the input data, essential for effective discrimination between real and fake data.
- LeakyReLU: Provides non-linearity to the learning process, helping to capture complex patterns in the data.

- Dense Layer: Outputs a scalar value that assesses the authenticity of the input data.

### 3. Classifier

- Purpose: Classifies both real and synthetic data into predefined categories.
- Architecture: Mirrors the discriminator in structure but concludes with a softmax layer to output class probabilities.
- Functionality: This component is crucial for ensuring that the synthetic data not only looks realistic but also adheres to the conditional distributions dependent on class labels.

## Training Mechanism

### Loss Functions

- Wasserstein Loss: Used for the discriminator to enhance training stability and convergence.
- Gradient Penalty: Implements the Lipschitz constraint, crucial for the Wasserstein GAN to function correctly.
- Categorical Crossentropy: Measures how well the classifier and the generator perform in predicting and generating correct labels, respectively.

### Optimizers

- Separate Adam optimizers for the generator, discriminator, and classifier with tailored learning rates and beta values to optimize each component effectively.

### Training Steps

- Noise Generation: Generate a noise vector for the generator to produce synthetic data.
- Discriminator Training: Update the discriminator by alternating between real and generated data to minimize its loss.
- Classifier Training: Train the classifier to maximize its accuracy on both real and generated data.
- Generator Training: Adjust the generator based on the feedback from both the discriminator and classifier, focusing on fooling the discriminator and accurately mimicking the class distributions.

## Parameter Influence and Experimentation

### Adjustable Parameters

- `input_dim` and `label_dim`: Define the shape and complexity of the data to be modeled. Adjusting these parameters allows the GAN to adapt to different sizes and types of data.
- `z_dim`: Controls the dimensionality of the latent space. A larger `z_dim` can capture more complex distributions but may lead to overfitting or mode collapse.
- `delta_mean` and `delta_sd`: Govern the strictness of privacy constraints by influencing the loss terms related to the distribution of features in real and synthetic data.

### Experimentation Opportunities

- **Varying Architecture:** Experiment with different architectures by adjusting the number and size of layers or by introducing new types of layers (e.g., dropout or different activation functions).
- **Hyperparameter Tuning:** Optimize learning rates, batch sizes, and training epochs to find the best settings for convergence and data quality.
- **Privacy Settings:** Modify `delta_mean` and `delta_sd` to explore trade-offs between data utility and privacy, adapting the model to various regulatory requirements.

### Usage Scenarios

- **Data Augmentation:** Enhance existing datasets with synthetic data to improve model training where real data is scarce or imbalanced.
- **Privacy-Preserving Data Sharing:** Generate synthetic datasets that can be shared externally without exposing sensitive information contained in the original data.